

# DuckDB & SQL

## Lecture 21

Dr. Colin Rundel

# SQL

Structures Query Language is a special purpose language for interacting with (querying and modifying) indexed tabular data.

- ANSI Standard but with dialect divergence (MySQL, Postgres, SQLite, etc.)
- This functionality maps very closely (but not exactly) with the data manipulation verbs present in dplyr.
- SQL is likely to be a foundational skill if you go into industry - learn it and put it on your CV

# DuckDB

DuckDB is an open-source column-oriented relational database management system (RDBMS) originally developed by Mark Raasveldt and Hannes Mühleisen at the Centrum Wiskunde & Informatica (CWI) in the Netherlands and first released in 2019. The project has over 6 million downloads per month. It is designed to provide high performance on complex queries against large databases in embedded configuration, such as combining tables with hundreds of columns and billions of rows. Unlike other embedded databases (for example, SQLite) DuckDB is not focusing on transactional (OLTP) applications and instead is specialized for online analytical processing (OLAP) workloads.

From [Wikipedia - DuckDB](#)

# DuckDB & DBI

DuckDB is a relational database just like SQLite and can be interacted with using DBI and the duckdb package.

```
1 library(DBI)
2 (con = dbConnect(duckdb::duckdb()))
```

```
<duckdb_connection c9240 driver=<duckdb_driver dbdir=':memory:' read_only=FALSE bigint
```

```
1 dbWriteTable(con, "flights", nycflights13::flights)
2 dbListTables(con)
```

```
[1] "flights"
```

```
1 dbGetQuery(con, "SELECT * FROM flights") |>
2   as_tibble()
```

```
# A tibble: 336,776 × 19
```

|    | year  | month | day   | dep_time | sched_dep_time | dep_delay | arr_time |
|----|-------|-------|-------|----------|----------------|-----------|----------|
|    | <int> | <int> | <int> | <int>    | <int>          | <dbl>     | <int>    |
| 1  | 2013  | 1     | 1     | 517      | 515            | 2         | 830      |
| 2  | 2013  | 1     | 1     | 533      | 529            | 4         | 850      |
| 3  | 2013  | 1     | 1     | 542      | 540            | 2         | 923      |
| 4  | 2013  | 1     | 1     | 544      | 545            | -1        | 1004     |
| 5  | 2013  | 1     | 1     | 554      | 600            | -6        | 812      |
| 6  | 2013  | 1     | 1     | 554      | 558            | -4        | 740      |
| 7  | 2013  | 1     | 1     | 555      | 600            | -5        | 913      |
| 8  | 2013  | 1     | 1     | 557      | 600            | -3        | 709      |
| 9  | 2013  | 1     | 1     | 557      | 600            | -3        | 838      |
| 10 | 2013  | 1     | 1     | 558      | 600            | -2        | 753      |

```
# i 336,766 more rows
```

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
# carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,  
# dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,  
# minute <dbl>, time_hour <dtm>
```

```

1 library(dplyr)
2 tbl(con, "flights") |>
3   filter(month == 10, day == 29) |>
4   count(origin, dest) |>
5   arrange(desc(n))

```

```

# Source:      SQL [?? x 3]
# Database:    DuckDB 1.5.0 [root@Darwin 25.3.0:R 4.5.3/:memory:]
# Ordered by: desc(n)
   origin dest      n
   <chr>  <chr> <dbl>
1 JFK     LAX      32
2 LGA     ORD      30
3 LGA     ATL      29
4 JFK     SFO      23
5 LGA     CLT      21
6 EWR     ORD      18
7 LGA     BOS      16
8 EWR     SFO      16
9 JFK     BOS      16
10 LGA    DCA      16
# i more rows

```

# DuckDB CLI

# Connecting via CLI

```
1 > #| eval: false  
2 > duckdb employees.duckdb
```

v1.2.1 8e52ec4395

Enter ".help" for usage hints.

```
1 D
```



# Table information

Dot commands are expressions that begins with `.` and are specific to the DuckDB CLI, some examples include:

```
1 D .tables
```

```
## employees
```

```
1 D .schema employees
```

```
## CREATE TABLE employees("name" VARCHAR, email VARCHAR, salary DOUBLE, dept VARCHAR);
```

```
1 D .indexes employees
```

```
1 D .maxrows 20
```

```
2 D .maxwidth 80
```

A full list of available dot commands can be found [here](#) or listed via `.help` in the CLI.

# SELECT Statements

1 D **SELECT** \* **FROM** employees;

##  
##  
##  
##  
##  
##  
##  
##  
##  
##  
##

| name<br>varchar | email<br>varchar  | salary<br>double | dept<br>varchar |
|-----------------|-------------------|------------------|-----------------|
| Alice           | alice@company.com | 52000.0          | Accounting      |
| Bob             | bob@company.com   | 40000.0          | Accounting      |
| Carol           | carol@company.com | 30000.0          | Sales           |
| Dave            | dave@company.com  | 33000.0          | Accounting      |
| Eve             | eve@company.com   | 44000.0          | Sales           |
| Frank           | frank@comany.com  | 37000.0          | Sales           |

# Output formats

The format of duckdb's output (in the CLI) is controlled via `.mode` - the default is `duckbox`, see possible [output formats](#).

```
1 D .mode csv
2 D SELECT * FROM employees;
```

```
## name,email,salary,dept
## Alice,alice@company.com,52000.0,Accounting
## Bob,bob@company.com,40000.0,Accounting
## Carol,carol@company.com,30000.0,Sales
## Dave,dave@company.com,33000.0,Accounting
## Eve,eve@company.com,44000.0,Sales
## Frank,frank@comany.com,37000.0,Sales
```

```
1 D .mode json
2 D SELECT * FROM employees;
```

```
## [{"name":"Alice","email":"alice@company.com","salary":52000.0
## {"name":"Bob","email":"bob@company.com","salary":40000.0,"dep
## {"name":"Carol","email":"carol@company.com","salary":30000.0,
## {"name":"Dave","email":"dave@company.com","salary":33000.0,"d
## {"name":"Eve","email":"eve@company.com","salary":44000.0,"dep
## {"name":"Frank","email":"frank@comany.com","salary":37000.0,"
```

```
1 D .mode markdown
2 D SELECT * FROM employees;
```

| ## | name  | email             | salary  | dept       |
|----|-------|-------------------|---------|------------|
| ## | Alice | alice@company.com | 52000.0 | Accounting |
| ## | Bob   | bob@company.com   | 40000.0 | Accounting |
| ## | Carol | carol@company.com | 30000.0 | Sales      |
| ## | Dave  | dave@company.com  | 33000.0 | Accounting |
| ## | Eve   | eve@company.com   | 44000.0 | Sales      |
| ## | Frank | frank@comany.com  | 37000.0 | Sales      |

```
1 D .mode insert
2 D SELECT * FROM employees;
```

```
INSERT INTO "table"("name",email,salary,dept) VALUES('Alice','al
INSERT INTO "table"("name",email,salary,dept) VALUES('Bob','bob@
INSERT INTO "table"("name",email,salary,dept) VALUES('Carol','ca
INSERT INTO "table"("name",email,salary,dept) VALUES('Dave','dav
INSERT INTO "table"("name",email,salary,dept) VALUES('Eve','eve@
INSERT INTO "table"("name",email,salary,dept) VALUES('Frank','fr
```

# A brief tour of SQL

# select() using SELECT

We can subset for certain columns (and rename them) using **SELECT**

```
1 D SELECT name AS first_name, salary FROM employees;
```

```
##
```

| first_name | salary  |
|------------|---------|
| varchar    | double  |
| Alice      | 52000.0 |
| Bob        | 40000.0 |
| Carol      | 30000.0 |
| Dave       | 33000.0 |
| Eve        | 44000.0 |
| Frank      | 37000.0 |

```
##
```

# arrange() using ORDER BY

We can sort our results by adding **ORDER BY** to our **SELECT** statement and reverse the ordering by include **DESC**.

```
1 D SELECT name AS first_name, salary FROM employees
2 D ORDER BY salary;
```

```
##
##
```

| first_name | salary  |
|------------|---------|
| varchar    | double  |
| Carol      | 30000.0 |
| Dave       | 33000.0 |
| Frank      | 37000.0 |
| Bob        | 40000.0 |
| Eve        | 44000.0 |
| Alice      | 52000.0 |

```
##
##
```

```
1 D SELECT name AS first_name, salary FROM employees
2 D ORDER BY salary DESC;
```

```
##
##
```

| first_name | salary  |
|------------|---------|
| varchar    | double  |
| Alice      | 52000.0 |
| Eve        | 44000.0 |
| Bob        | 40000.0 |
| Frank      | 37000.0 |
| Dave       | 33000.0 |
| Carol      | 30000.0 |

```
##
##
```

# filter() using WHERE

We can filter rows using a **WHERE** clause

```
1 D SELECT * FROM employees WHERE salary < 40000;
```

##

| name<br>varchar | email<br>varchar  | salary<br>double | dept<br>varchar |
|-----------------|-------------------|------------------|-----------------|
| Carol           | carol@company.com | 30000.0          | Sales           |
| Dave            | dave@company.com  | 33000.0          | Accounting      |
| Frank           | frank@comany.com  | 37000.0          | Sales           |

##

```
1 D SELECT * FROM employees WHERE salary < 40000 AND dept = 'Sales';
```

##

| name<br>varchar | email<br>varchar  | salary<br>double | dept<br>varchar |
|-----------------|-------------------|------------------|-----------------|
| Carol           | carol@company.com | 30000.0          | Sales           |
| Frank           | frank@comany.com  | 37000.0          | Sales           |

##

# group\_by() and summarize() via GROUP BY

We can create groups for the purpose of summarizing using **GROUP BY**.

```
1 D SELECT dept, COUNT(*) AS n FROM employees GROUP BY dept;
```

```
##  
##  
##  
##  
##  
##  
##
```

| dept<br>varchar | n<br>int64 |
|-----------------|------------|
| Sales           | 3          |
| Accounting      | 3          |



# head() using LIMIT

We can limit the number of results we get by using **LIMIT**

```
1 D SELECT * FROM employees LIMIT 3;
```

```
##
```

| name<br>varchar | email<br>varchar  | salary<br>double | dept<br>varchar |
|-----------------|-------------------|------------------|-----------------|
| Alice           | alice@company.com | 52000.0          | Accounting      |
| Bob             | bob@company.com   | 40000.0          | Accounting      |
| Carol           | carol@company.com | 30000.0          | Sales           |

```
##
```

# Exercise 1

Using duckdb calculate the following quantities for `employees.duckdb`,

1. The total costs in payroll for this company
2. The average salary within each department

# Reading from CSV files

DuckDB has a neat trick in that it can treat files as tables (for supported formats), this lets you query them without having to explicitly create a table in the database and .

We can also make this explicit by using the `read_csv()` function, which is useful if we need to use custom options (e.g. specify a different delimiter)

```
1 D SELECT * FROM 'phone.csv';
```

|    |         |              |
|----|---------|--------------|
| ## |         |              |
| ## | name    | phone        |
| ## | varchar | varchar      |
| ## |         |              |
| ## | Bob     | 919 555-1111 |
| ## | Carol   | 919 555-2222 |
| ## | Eve     | 919 555-3333 |
| ## | Frank   | 919 555-4444 |
| ## |         |              |

```
1 D SELECT * FROM
2 D   read_csv('phone.csv', delim = ',');
```

|    |         |              |
|----|---------|--------------|
| ## |         |              |
| ## | name    | phone        |
| ## | varchar | varchar      |
| ## |         |              |
| ## | Bob     | 919 555-1111 |
| ## | Carol   | 919 555-2222 |
| ## | Eve     | 919 555-3333 |
| ## | Frank   | 919 555-4444 |
| ## |         |              |

# Tables from CSV

If we wanted to explicitly create a table from the CSV file this is also possible,

```
1 D .tables
```

## employees

```
1 D CREATE TABLE phone AS
2 D SELECT * FROM 'phone.csv';
```

```
1 D .tables
```

## employees phone

```
1 D SELECT * FROM phone;
```

##

| name<br>varchar | phone<br>varchar |
|-----------------|------------------|
| Bob             | 919 555-1111     |
| Carol           | 919 555-2222     |
| Eve             | 919 555-3333     |
| Frank           | 919 555-4444     |

##

# Views from CSV

It is also possible to create a view from a file - this acts like a table but the data is not copied from the file.

```
1 D .tables
```

## employees

```
1 D CREATE VIEW phone_view AS
2 D   SELECT * FROM 'phone.csv';
```

```
1 D .tables
```

## employees phone phone\_view

```
1 D SELECT * FROM phone_view;
```

##

| ## | name    | phone        |
|----|---------|--------------|
| ## | varchar | varchar      |
| ## | Bob     | 919 555-1111 |
| ## | Carol   | 919 555-2222 |
| ## | Eve     | 919 555-3333 |
| ## | Frank   | 919 555-4444 |
| ## |         |              |

# Deleting tables and views

Tables and views can be deleted using **DROP**

```
1 D .tables
```

```
## employees  phone  phone_view
```

```
1 D DROP TABLE phone;  
2 D DROP VIEW phone_view;
```

```
1 D .tables
```

```
## employees
```

# Joins - Default

If not otherwise specified the default join in DuckDB will be an inner join.

```
1 D SELECT * FROM employees JOIN phone;
```

```
## Parser Error: syntax error at or near ";"  
## LINE 1: SELECT * FROM employees JOIN phone;
```

this error occurs because duckdb requires an **ON** or **USING** clause unless using **NATURAL**.

```
1 D SELECT * FROM employees NATURAL JOIN phone;
```

```
##  
##  
##  
##  
##  
##  
##  
##  
##
```

| name<br>varchar | email<br>varchar  | salary<br>double | dept<br>varchar | phone<br>varchar |
|-----------------|-------------------|------------------|-----------------|------------------|
| Bob             | bob@company.com   | 40000.0          | Accounting      | 919 555-1111     |
| Carol           | carol@company.com | 30000.0          | Sales           | 919 555-2222     |
| Eve             | eve@company.com   | 44000.0          | Sales           | 919 555-3333     |
| Frank           | frank@comany.com  | 37000.0          | Sales           | 919 555-4444     |

# Inner Join - Explicit

```
1 D SELECT * FROM employees JOIN phone ON employees.name = phone.name;
```

|    |         |                   |         |            |         |              |
|----|---------|-------------------|---------|------------|---------|--------------|
| ## | name    | email             | salary  | dept       | name    | phone        |
| ## | varchar | varchar           | double  | varchar    | varchar | varchar      |
| ## | Bob     | bob@company.com   | 40000.0 | Accounting | Bob     | 919 555-1111 |
| ## | Carol   | carol@company.com | 30000.0 | Sales      | Carol   | 919 555-2222 |
| ## | Eve     | eve@company.com   | 44000.0 | Sales      | Eve     | 919 555-3333 |
| ## | Frank   | frank@comany.com  | 37000.0 | Sales      | Frank   | 919 555-4444 |
| ## |         |                   |         |            |         | ...          |

to avoid the duplicate `name` column we can specify `USING` instead of `ON`

```
1 D SELECT * FROM employees JOIN phone USING(name);
```

|    |         |                   |         |            |              |
|----|---------|-------------------|---------|------------|--------------|
| ## | name    | email             | salary  | dept       | phone        |
| ## | varchar | varchar           | double  | varchar    | varchar      |
| ## | Bob     | bob@company.com   | 40000.0 | Accounting | 919 555-1111 |
| ## | Carol   | carol@company.com | 30000.0 | Sales      | 919 555-2222 |
| ## | Eve     | eve@company.com   | 44000.0 | Sales      | 919 555-3333 |
| ## | Frank   | frank@comany.com  | 37000.0 | Sales      | 919 555-4444 |
| ## |         |                   |         |            |              |



# Left Join - Natural

```
1 D SELECT * FROM employees NATURAL LEFT JOIN phone;
```

|    |         |                   |         |            |              |
|----|---------|-------------------|---------|------------|--------------|
| ## | name    | email             | salary  | dept       | phone        |
| ## | varchar | varchar           | double  | varchar    | varchar      |
| ## | Bob     | bob@company.com   | 40000.0 | Accounting | 919 555-1111 |
| ## | Carol   | carol@company.com | 30000.0 | Sales      | 919 555-2222 |
| ## | Eve     | eve@company.com   | 44000.0 | Sales      | 919 555-3333 |
| ## | Frank   | frank@comany.com  | 37000.0 | Sales      | 919 555-4444 |
| ## | Alice   | alice@company.com | 52000.0 | Accounting |              |
| ## | Dave    | dave@company.com  | 33000.0 | Accounting |              |
| ## |         |                   |         |            |              |

# Left Join - Explicit

```
1 D SELECT * FROM employees LEFT JOIN phone ON employees.name = phone.name;
```

|    |         |                   |         |            |         |              |
|----|---------|-------------------|---------|------------|---------|--------------|
| ## | name    | email             | salary  | dept       | name    | phone        |
| ## | varchar | varchar           | double  | varchar    | varchar | varchar      |
| ## | Bob     | bob@company.com   | 40000.0 | Accounting | Bob     | 919 555-1111 |
| ## | Carol   | carol@company.com | 30000.0 | Sales      | Carol   | 919 555-2222 |
| ## | Eve     | eve@company.com   | 44000.0 | Sales      | Eve     | 919 555-3333 |
| ## | Frank   | frank@comany.com  | 37000.0 | Sales      | Frank   | 919 555-4444 |
| ## | Alice   | alice@company.com | 52000.0 | Accounting |         |              |
| ## | Dave    | dave@company.com  | 33000.0 | Accounting |         |              |
| ## |         |                   |         |            |         |              |

Duplicate **name** column can be avoided by more restrictive **SELECT**,

```
1 D SELECT employees.*, phone FROM employees LEFT JOIN phone ON employees.name = phone.name;
```

|    |         |                   |         |            |              |
|----|---------|-------------------|---------|------------|--------------|
| ## | name    | email             | salary  | dept       | phone        |
| ## | varchar | varchar           | double  | varchar    | varchar      |
| ## | Bob     | bob@company.com   | 40000.0 | Accounting | 919 555-1111 |
| ## | Carol   | carol@company.com | 30000.0 | Sales      | 919 555-2222 |
| ## | Eve     | eve@company.com   | 44000.0 | Sales      | 919 555-3333 |
| ## | Frank   | frank@comany.com  | 37000.0 | Sales      | 919 555-4444 |
| ## | Alice   | alice@company.com | 52000.0 | Accounting |              |
| ## | Dave    | dave@company.com  | 33000.0 | Accounting |              |
| ## |         |                   |         |            |              |

# Other Joins

As you would expect all other standard joins are supported including **RIGHT JOIN**, **FULL JOIN**, **CROSS JOIN**, **SEMI JOIN**, **ANTI JOIN**, etc.

```
1 D SELECT employees.*, phone FROM employees NATURAL FULL JOIN phone;
```

##

| name<br>varchar | email<br>varchar  | salary<br>double | dept<br>varchar | phone<br>varchar |
|-----------------|-------------------|------------------|-----------------|------------------|
| Bob             | bob@company.com   | 40000.0          | Accounting      | 919 555-1111     |
| Carol           | carol@company.com | 30000.0          | Sales           | 919 555-2222     |
| Eve             | eve@company.com   | 44000.0          | Sales           | 919 555-3333     |
| Frank           | frank@comany.com  | 37000.0          | Sales           | 919 555-4444     |
| Alice           | alice@company.com | 52000.0          | Accounting      |                  |
| Dave            | dave@company.com  | 33000.0          | Accounting      |                  |

##

```
1 D SELECT employees.*, phone FROM employees NATURAL RIGHT JOIN phone;
```

##

| name<br>varchar | email<br>varchar  | salary<br>double | dept<br>varchar | phone<br>varchar |
|-----------------|-------------------|------------------|-----------------|------------------|
| Bob             | bob@company.com   | 40000.0          | Accounting      | 919 555-1111     |
| Carol           | carol@company.com | 30000.0          | Sales           | 919 555-2222     |
| Eve             | eve@company.com   | 44000.0          | Sales           | 919 555-3333     |
| Frank           | frank@comany.com  | 37000.0          | Sales           | 919 555-4444     |

##

# Subqueries

Tables can be nested within tables for the purpose of queries,

```
1 D SELECT * FROM (  
2 D   SELECT * FROM employees NATURAL LEFT JOIN phone  
3 D ) combined WHERE phone IS NULL;
```

##

| name<br>varchar | email<br>varchar  | salary<br>double | dept<br>varchar | phone<br>varchar |
|-----------------|-------------------|------------------|-----------------|------------------|
| Alice           | alice@company.com | 52000.0          | Accounting      |                  |
| Dave            | dave@company.com  | 33000.0          | Accounting      |                  |

##

```
1 D SELECT * FROM (  
2 D   SELECT * FROM employees NATURAL LEFT JOIN phone  
3 D ) combined WHERE phone IS NOT NULL;
```

##

| name<br>varchar | email<br>varchar  | salary<br>double | dept<br>varchar | phone<br>varchar |
|-----------------|-------------------|------------------|-----------------|------------------|
| Bob             | bob@company.com   | 40000.0          | Accounting      | 919 555-1111     |
| Carol           | carol@company.com | 30000.0          | Sales           | 919 555-2222     |
| Eve             | eve@company.com   | 44000.0          | Sales           | 919 555-3333     |
| Frank           | frank@comany.com  | 37000.0          | Sales           | 919 555-4444     |

##

## Exercise 2

Lets try to create a table that has a new column - `abv_avg` which contains how much more (or less) than the average, for their department, each person is paid.

Hint - This will require joining a subquery.

# Query plan

# Setup

To give us a bit more variety (and data), we have created another SQLite database `flights.duckdb` that contains both `nycflights13::flights` and `nycflights13::planes`, the latter of which has details on the characteristics of the planes in the dataset as identified by their tail numbers.

```
1 db = DBI::dbConnect(duckdb::duckdb(), "flights.duckdb")
2 dplyr::copy_to(db, nycflights13::flights, name = "flights", temporary = FALSE, over
3 dplyr::copy_to(db, nycflights13::planes, name = "planes", temporary = FALSE, overv
4 DBI::dbDisconnect(db)
```

All of the following code will be run in the DuckDB command line interface, make sure you've created the database and copied both the flights and planes tables into the db or use the version provided in the [exercises/](#) repo.

# Opening `flights.duckdb`

The database can then be opened from the terminal tab using,

```
1 > #| eval: false  
2 > duckdb flights.duckdb
```

As before we set a couple of configuration options so that our output is readable. We also include `.timer on` so that we get timings for our queries.

```
1 D .maxrows 20  
2 D .maxwidth 80  
3 D .timer on
```



# flights

```
1 D SELECT * FROM flights LIMIT 10;
```

|    |                                                     |       |       |     |          |        |        |                      |
|----|-----------------------------------------------------|-------|-------|-----|----------|--------|--------|----------------------|
| ## | ##                                                  | ##    | ##    | ##  | ##       | ##     | ##     | ##                   |
| ## | year                                                | month | day   | ... | distance | hour   | minute | time_hour            |
| ## | int32                                               | int32 | int32 |     | double   | double | double | timestamp            |
| ## | 2013                                                | 1     | 1     | ... | 1400.0   | 5.0    | 15.0   | 2013-01-01 10:00:00  |
| ## | 2013                                                | 1     | 1     | ... | 1416.0   | 5.0    | 29.0   | 2013-01-01 10:00:00  |
| ## | 2013                                                | 1     | 1     | ... | 1089.0   | 5.0    | 40.0   | 2013-01-01 10:00:00  |
| ## | 2013                                                | 1     | 1     | ... | 1576.0   | 5.0    | 45.0   | 2013-01-01 10:00:00  |
| ## | 2013                                                | 1     | 1     | ... | 762.0    | 6.0    | 0.0    | 2013-01-01 11:00:00  |
| ## | 2013                                                | 1     | 1     | ... | 719.0    | 5.0    | 58.0   | 2013-01-01 10:00:00  |
| ## | 2013                                                | 1     | 1     | ... | 1065.0   | 6.0    | 0.0    | 2013-01-01 11:00:00  |
| ## | 2013                                                | 1     | 1     | ... | 229.0    | 6.0    | 0.0    | 2013-01-01 11:00:00  |
| ## | 2013                                                | 1     | 1     | ... | 944.0    | 6.0    | 0.0    | 2013-01-01 11:00:00  |
| ## | 2013                                                | 1     | 1     | ... | 733.0    | 6.0    | 0.0    | 2013-01-01 11:00:00  |
| ## | 10 rows                                             |       |       |     |          |        |        | 19 columns (7 shown) |
| ## | Run Time (s): real 0.020 user 0.000784 sys 0.002284 |       |       |     |          |        |        |                      |

# planes

```
1 D SELECT * FROM planes LIMIT 10;
```

##

| tailnum<br>varchar | year<br>int32 | type<br>varchar       | ... | seats<br>int32 | speed<br>int32      | engine<br>varchar |
|--------------------|---------------|-----------------------|-----|----------------|---------------------|-------------------|
| N10156             | 2004          | Fixed wing multi e... | ... | 55             |                     | Turbo-fan         |
| N102UW             | 1998          | Fixed wing multi e... | ... | 182            |                     | Turbo-fan         |
| N103US             | 1999          | Fixed wing multi e... | ... | 182            |                     | Turbo-fan         |
| N104UW             | 1999          | Fixed wing multi e... | ... | 182            |                     | Turbo-fan         |
| N10575             | 2002          | Fixed wing multi e... | ... | 55             |                     | Turbo-fan         |
| N105UW             | 1999          | Fixed wing multi e... | ... | 182            |                     | Turbo-fan         |
| N107US             | 1999          | Fixed wing multi e... | ... | 182            |                     | Turbo-fan         |
| N108UW             | 1999          | Fixed wing multi e... | ... | 182            |                     | Turbo-fan         |
| N109UW             | 1999          | Fixed wing multi e... | ... | 182            |                     | Turbo-fan         |
| N110UW             | 1999          | Fixed wing multi e... | ... | 182            |                     | Turbo-fan         |
| 10 rows            |               |                       |     |                | 9 columns (6 shown) |                   |

##

## Run Time (s): real 0.003 user 0.000819 sys 0.000018

## Exercise 3

Write a query that determines the total number of seats available on all of the planes that flew out of New York in 2013.

# A Solution?

Does the following seem correct?

```
1 D SELECT sum(seats) FROM flights NATURAL LEFT JOIN planes;
```

```
##  
##  
##  
##  
##  
##
```

|            |
|------------|
| sum(seats) |
| int128     |
| 614366     |

```
## Run Time (s): real 0.012 user 0.016061 sys 0.002386
```

Why?

# Correct solution

Join and select:

```
1 D SELECT sum(seats) FROM flights LEFT JOIN planes USING (tailnum);
```


##

|    |            |
|----|------------|
| ## | sum(seats) |
| ## | int128     |
| ## | 38851317   |
| ## |            |

## Run Time (s): real 0.005 user 0.010150 sys 0.000291

# EXPLAIN

```
1 D EXPLAIN SELECT sum(seats) FROM flights LEFT JOIN planes USING (tailnum);
```




# EXPLAIN ANALYZE

```
1 D EXPLAIN ANALYZE SELECT sum(seats) FROM flights LEFT JOIN planes USING (tailnum);
```

```
##
##
##  Query Profiling Information
##
##
## EXPLAIN ANALYZE SELECT sum(seats) FROM flights LEFT JOIN planes USING (tailnum);
```

Total Time: 0.0045s



# dplyr

```
1 library(dplyr)
2 flights = nycflights13::flights
3 planes = nycflights13::planes
4
5 system.time({
6   flights |>
7     left_join(planes, by = c("tailnum" = "tailnum")) |>
8     summarise(total_seats = sum(seats, na.rm = TRUE))
9 })
```

| user  | system | elapsed |
|-------|--------|---------|
| 0.050 | 0.004  | 0.054   |



# NYC Taxi Demo